

MATLAB Distributed Compute on the Cray CS300

Sean Cleveland Ph.D, Ron Merrill Ph.D,
David Schanzenbach M.S.

Information Technology Services
Department of Cyberinfrastructure
University of Hawai'i

<https://www.hawaii.edu/its/ci/>
uh-hpc-help@lists.hawaii.edu

January 18, 2016



UNIVERSITY OF HAWAII

Overview

- 1 MATLAB Distributed Compute
- 2 Configuration
 - Setting up the MATLAB path
 - Importing the Cray CS300 MATLAB profile
- 3 Basic Commands & Testing
 - Commands
 - Validation
- 4 Demo
- 5 Resources



UNIVERSITY OF HAWAII

Outline

- 1 MATLAB Distributed Compute
- 2 Configuration
 - Setting up the MATLAB path
 - Importing the Cray CS300 MATLAB profile
- 3 Basic Commands & Testing
 - Commands
 - Validation
- 4 Demo
- 5 Resources



UNIVERSITY OF HAWAII

MATLAB on the Cluster

- MATLAB Distributed Compute Server (DCS) R2015b is installed on the Cray CS300
- We currently hold a license for 64 *Workers*.
 - One worker is equal to an instance of MATLAB
- In order to use MATLAB on the Cray CS300, you will need to own the MATLAB product with the Parallel Computing Toolbox
- The workers on the Cray CS300 contain all the MATLAB toolboxes
 - **This does not mean you get free access to all toolboxes**
 - **You will only have access to the MATLAB toolboxes you have purchased**



Pre-requisites to using MATLAB on the Cray CS300

In order to use MATLAB on the cluster you must meet the follow requirements

- ① Own a copy of MATLAB, with a license for the Parallel Computing Toolbox
- ② Have gone through the On-boarding course for the Cray CS300



Outline

- 1 MATLAB Distributed Compute
- 2 Configuration
 - Setting up the MATLAB path
 - Importing the Cray CS300 MATLAB profile
- 3 Basic Commands & Testing
 - Commands
 - Validation
- 4 Demo
- 5 Resources



UNIVERSITY OF HAWAII

Required Files

In order to submit MATLAB jobs to the Cray CS300, a packet of several scripts, and a configuration file are required:

- Download the zip file containing the MATLAB setup files here:
<http://go.hawaii.edu/YH>
- Download should contain:
 - A directory named `slurm_matlab`, which contains several `.m` and `.sh` files
 - A settings file named `uhhpc.settings`



Setting up MATLAB

- Unzip setup files
- Move the *slurm_matlab* directory to a location it can reside permanently
 - This directory will be added to your MATLAB path
- Start up MATLAB

The next few Slides will cover what needs to be done in MATLAB to allow communication with the Cray CS300 from MATLAB



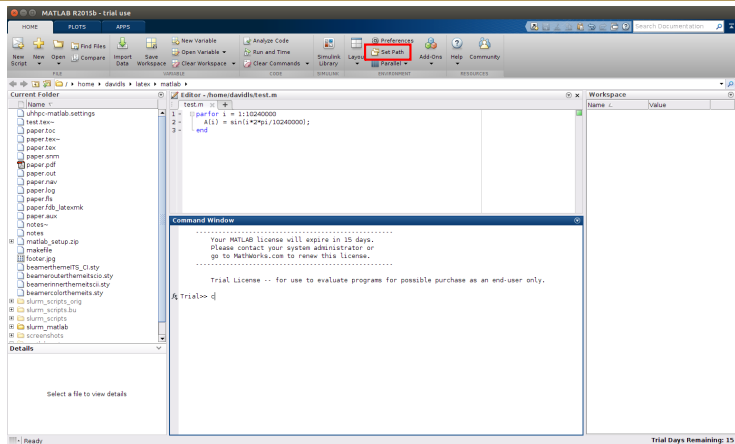
Outline

- 1 MATLAB Distributed Compute
- 2 Configuration
 - Setting up the MATLAB path
 - Importing the Cray CS300 MATLAB profile
- 3 Basic Commands & Testing
 - Commands
 - Validation
- 4 Demo
- 5 Resources



UNIVERSITY OF HAWAII

Adding to the MATLAB Path

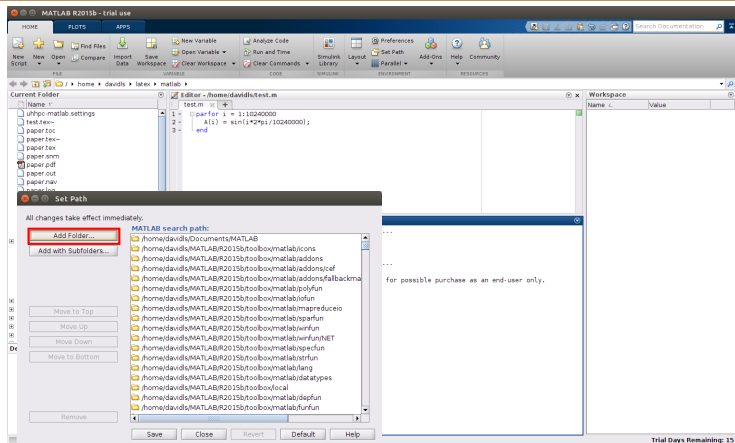


Click the “**Set Path**” button under the “**Home**” tab in the MATLAB UI



UNIVERSITY OF HAWAII

Adding to the MATLAB Path

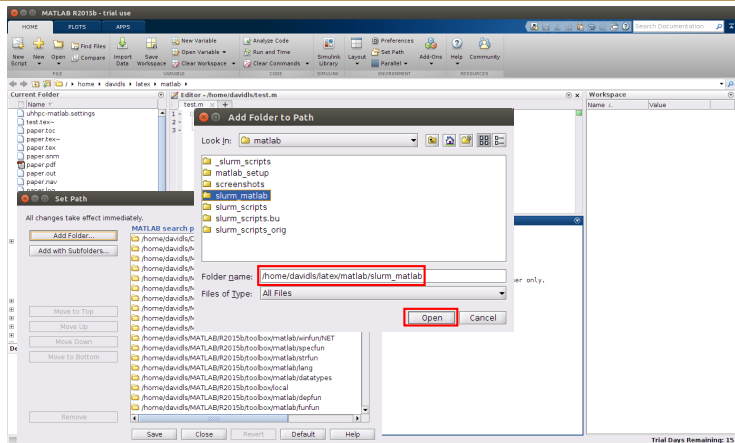


In the Set Path dialogue, click **“Add Folder ...”**



UNIVERSITY OF HAWAII

Adding to the MATLAB Path

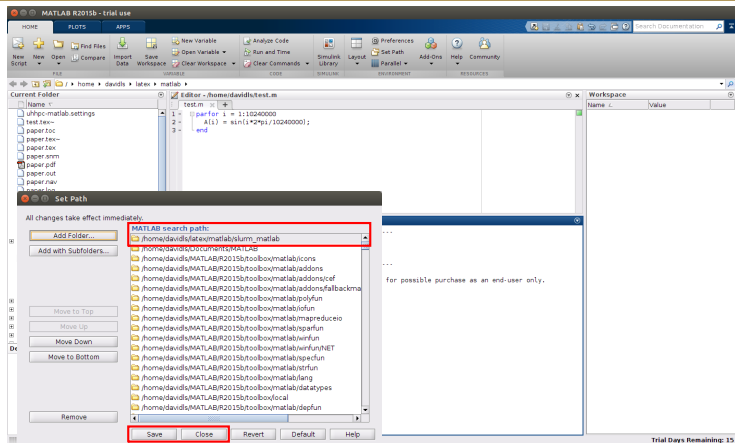


In the Add Folder to Path dialogue, navigate to the directory the **slurm_matlab** folder resides in and select and **Open** it



UNIVERSITY OF HAWAII

Adding to the MATLAB Path



The newly added path should appear in the “**MATLAB** serach path:” list. Once it added, go ahead and click “**Save**”, then “**Close**”



UNIVERSITY OF HAWAI'I

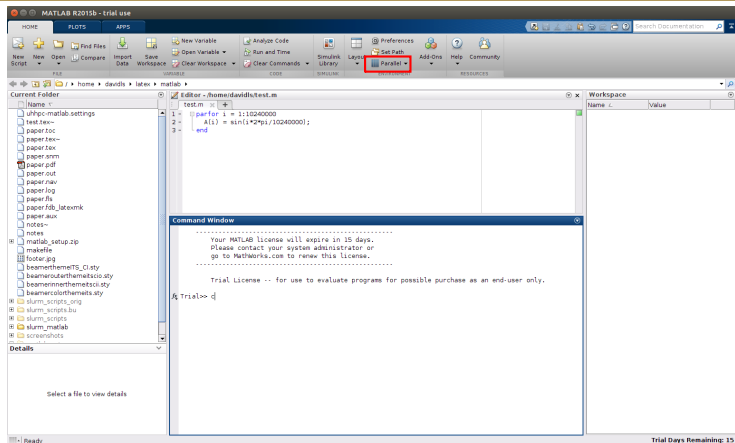
Outline

- 1 MATLAB Distributed Compute
- 2 Configuration
 - Setting up the MATLAB path
 - Importing the Cray CS300 MATLAB profile
- 3 Basic Commands & Testing
 - Commands
 - Validation
- 4 Demo
- 5 Resources



UNIVERSITY OF HAWAII

Cluster Profile Setup

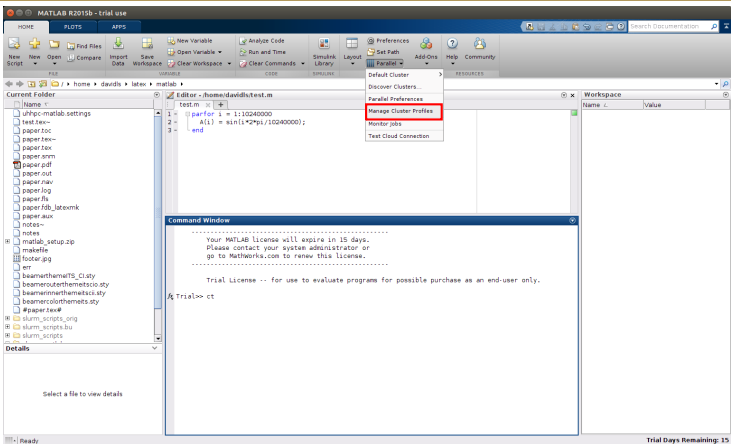


Click the “**Parallel**” button to bring up the Parallel drop down menu



UNIVERSITY OF HAWAII

Cluster Profile Setup

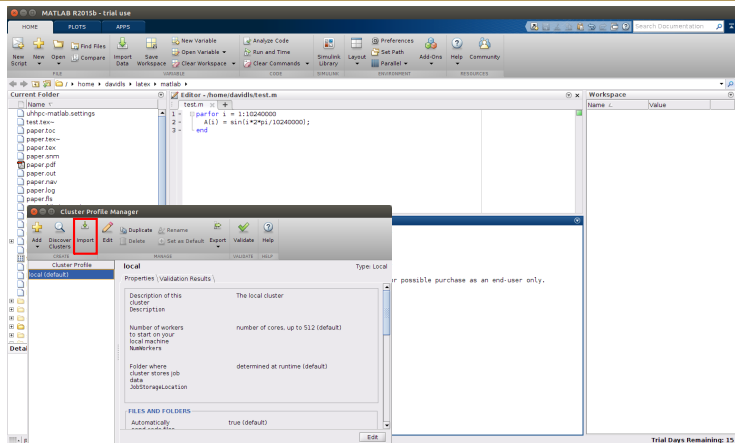


Select **Manage Cluster Profiles** in the Parallel drop down menu



UNIVERSITY OF HAWAII

Cluster Profile Setup

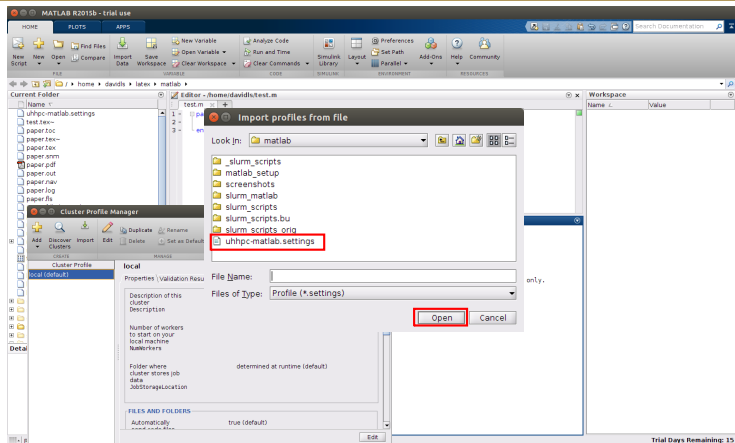


In the **Cluster Profile Manager** dialogue box, click on the **“Import”** icon



UNIVERSITY OF HAWAI'I

Cluster Profile Setup

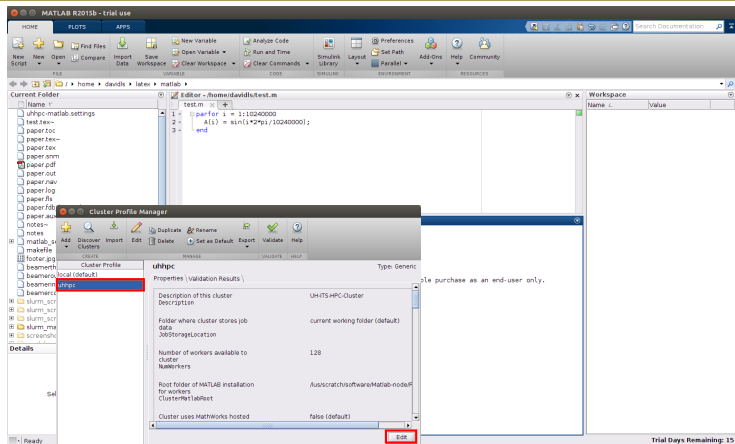


Navigate to and **Open** the uhhpc-matlab.settings file that was downloaded



UNIVERSITY OF HAWAII

Cluster Profile Setup

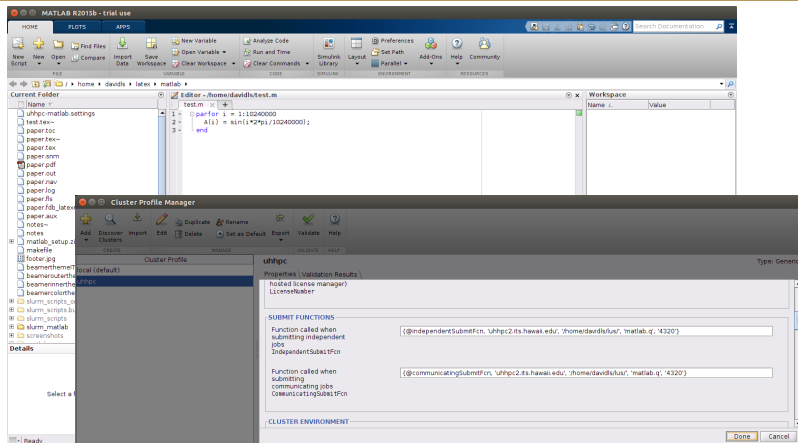


Select the **uhhpc** cluster profile, and press the “**Edit**” button.
For ease of editing, you may want to increase the window size



UNIVERSITY OF HAWAI'I

Cluster Profile Setup




Each of the functions in the **SUBMIT FUNCTIONS** section, have the follow format:

{ '<func>', '<server addr>', '<remote storage dir>', '<slurm partition>', '<slurm timelimit>' }



UNIVERSITY OF HAWAII

Cluster Profile Setup



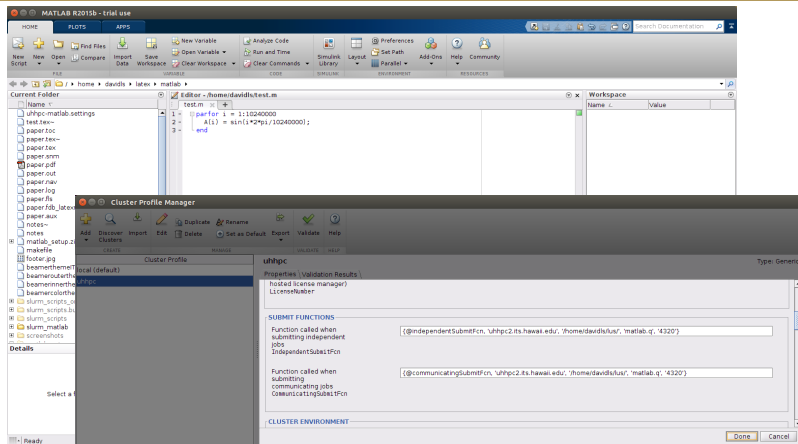
{ '<func>', '<server addr>', '<remote storage dir>', '<slurm partition>', '<slurm timelimit>' }

Replace the remote storage dir with the path that matches your users
lus folder location i.e., /home/<username>/lus/



UNIVERSITY OF HAWAII

Cluster Profile Setup



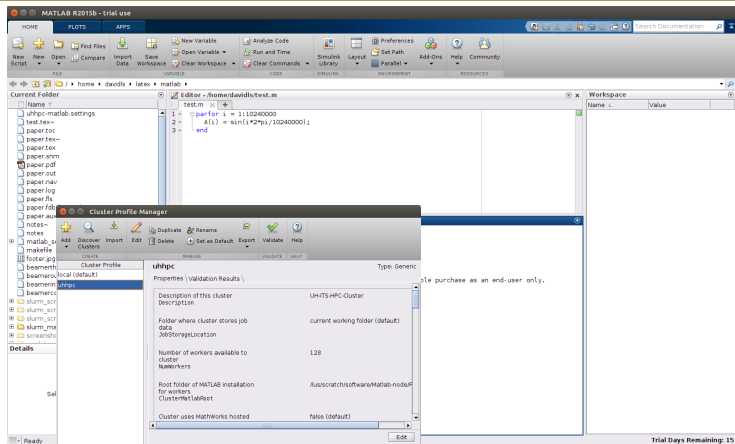
{ '<func>', '<server addr>', '<remote storage dir>', '<slurm partition>', '<slurm timelimit>' }

Adjust the slurm partition and slurm timelimit to fit the partition you wish to run on, and finally click **done**



UNIVERSITY OF HAWAII

Cluster Profile Setup



The profile should now be configured for your user on the Cray CS300



UNIVERSITY OF HAWAI'I

Cluster Profile

- To simplify the use of MATLAB on the cluster, we suggest setting the UHHPC cluster profile as the default profile
 - Many of the parallel toolbox commands will use the default cluster when an alternate cluster object is not provided
- If the UHHPC cluster profile is not set as the default, the cluster profile name, '**uhhpc**', must be provided when creating the cluster object
- If you plan to change partitions or run in multiple partitions on the Cray CS300, you will need to modify the existing profile, or create duplicate profiles with the required changes to the remote storage location, partition and timelimit



Outline

- 1 MATLAB Distributed Compute
- 2 Configuration
 - Setting up the MATLAB path
 - Importing the Cray CS300 MATLAB profile
- 3 Basic Commands & Testing
 - Commands
 - Validation
- 4 Demo
- 5 Resources



UNIVERSITY OF HAWAII

Outline

- 1 MATLAB Distributed Compute
- 2 Configuration
 - Setting up the MATLAB path
 - Importing the Cray CS300 MATLAB profile
- 3 Basic Commands & Testing
 - **Commands**
 - Validation
- 4 Demo
- 5 Resources



UNIVERSITY OF HAWAII

- **pctconfig** – Configure settings for Parallel Computing Toolbox client session
 - Is not persistent
 - Must be set before any interaction with the parallel toolbox occurs
 - Used to set the correct 'hostname' in MATLAB
 - e.g., `> pctconfig('hostname', '128.171.173.227')`
 - A valid 'hostname' is required for validation to pass as well as for many of the parallel functions to work
- **parcluster** – Creates a cluster object which can be passed to functions used to create parallel tasks
- **gcp** – Get current parallel pool
- **addAttachedFiles** – Attach files or folders to parallel pool



Commands

- **batch** – Run MATLAB script or function on worker
 - Non-interactive. MATLAB can be closed
- **parpool** – Create parallel pool on cluster
- **spmd** – Execute code in parallel on workers of parallel pool

For more information and other commands, visit the **Distributed toolbox help manual** on the MathWorks website



Outline

- 1 MATLAB Distributed Compute
- 2 Configuration
 - Setting up the MATLAB path
 - Importing the Cray CS300 MATLAB profile
- 3 Basic Commands & Testing
 - Commands
 - **Validation**
- 4 Demo
- 5 Resources



UNIVERSITY OF HAWAII

Simple Tests

- Simplest test is to use “*validate*” in the **Cluster Profile Manager**
- Before testng, make sure to reduce the value for “Number of workers available to cluster” (NumWorkers) value to something lower than 64
 - Typically we test with 2 – 10 workers
- NumWorkers must be lowered since validate uses this as the number of workers to request
- You also must make sure the hostname in MATLAB is set correctly (using pctconf), or the parpool test may fail
- Remember to reset the NumWorkers value back to 64 or you may get errors saying not enough workers when requesting more than the current NumWorkers value



Outline

- 1 MATLAB Distributed Compute
- 2 Configuration
 - Setting up the MATLAB path
 - Importing the Cray CS300 MATLAB profile
- 3 Basic Commands & Testing
 - Commands
 - Validation
- 4 Demo
- 5 Resources



UNIVERSITY OF HAWAII

Live Demo ...



Outline

- 1 MATLAB Distributed Compute
- 2 Configuration
 - Setting up the MATLAB path
 - Importing the Cray CS300 MATLAB profile
- 3 Basic Commands & Testing
 - Commands
 - Validation
- 4 Demo
- 5 Resources



UNIVERSITY OF HAWAII

- Newest revision of these slides – <http://go.hawaii.edu/H5>
- Parallel MATLAB (PCT and DCS) @ Harvard
 - Not all sections pertain to our cluster, but some sections provide some information on running jobs with the parallel toolbox
 - The “Submitting DCS jobs from within MATLAB” section covers using the batch command to submit parallel jobs to a cluster



Questions?



UNIVERSITY OF HAWAII